

# Real Estate Regression

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import plotly.express as px
        4 import matplotlib.pyplot as plt
        5 import seaborn as sb
        6 import folium
        7 import joblib
        8 import warnings
        9
       10 from sklearn.impute import SimpleImputer
       11 from sklearn.linear_model import LinearRegression, Ridge
       12 from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
       13 from sklearn.metrics import mean_squared_error, r2_score
       14 from sklearn.pipeline import make_pipeline, Pipeline
       15 from sklearn.model_selection import train_test_split
       16 from sklearn.compose import ColumnTransformer
       17 from sklearn.preprocessing import StandardScaler
       18
       19
       20 %matplotlib inline
       21 sb.set_theme(style = 'white')
       22 warnings.filterwarnings('ignore')
```

```

In [2]: 1 # plot functions
        2 # User defined histogram plot function
        3 def hist_plot(data, x_arg, title, x_label, y_label, bin_size, kde):
        4     sb.histplot(data = data, x = x_arg, bins = bin_size, kde = kde)
        5     plt.axvline(data[x_arg].mean(), color = 'red', linestyle = '-', linewidth = 2, label = 'Mean')
        6     plt.axvline(data[x_arg].median(), color = 'cyan', linestyle = ':', linewidth = 2, label = 'Median')
        7     plt.title(title, size = 15, weight = 'bold')
        8     plt.xlabel(x_label, size = 15, weight = 'bold')
        9     plt.ylabel(y_label, size = 15, weight = 'bold')
       10     plt.legend()
       11
       12 # User defined box plot function
       13 def box_plot(data, x_arg, y_arg, title, x_label):
       14     ax = sb.boxplot(data = data, x = x_arg, y = y_arg, color = 'white')#sb.color_palette()[0])
       15     plt.title(title, size = 15, weight = 'bold')
       16     plt.xlabel(x_label, size = 15, weight = 'bold')
       17     plt.setp(ax.lines, color = 'black')
       18
       19 def scatter_plot(data, x_arg, y_arg, x_label, y_label, title):
       20     sb.scatterplot(data = data, x = x_arg, y = y_arg)
       21     plt.title(title, size = 12, weight = 'bold')
       22     plt.xlabel(x_label, size = 10, weight = 'bold')
       23     plt.ylabel(y_label, size = 10, weight = 'bold')
       24     plt.show()

```

```

In [3]: 1 re_data = pd.read_csv('real_estate.csv')
        2 re_data.head()

```

```

Out[3]:

```

	transaction_date	house_age	transit_distance	local_convenience_stores	latitude	longitude	price_per_unit
0	2012.917	32.0	84.87882	10	24.98298	121.54024	37.9
1	2012.917	19.5	306.59470	9	24.98034	121.53951	42.2
2	2013.583	13.3	561.98450	5	24.98746	121.54391	47.3
3	2013.500	13.3	561.98450	5	24.98746	121.54391	54.8
4	2012.833	5.0	390.56840	5	24.97937	121.54245	43.1

# Exploratory Analysis

In [4]: 1 re\_data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 414 entries, 0 to 413
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   transaction_date       414 non-null   float64
1   house_age              414 non-null   float64
2   transit_distance       414 non-null   float64
3   local_convenience_stores 414 non-null   int64
4   latitude               414 non-null   float64
5   longitude              414 non-null   float64
6   price_per_unit         414 non-null   float64
dtypes: float64(6), int64(1)
memory usage: 22.8 KB
```

Form the results, the dataset is complete and their are no missing values.

## Estate location

In [5]:

```
1  fig = px.scatter_mapbox(re_data,
2                          lat = 'latitude',
3                          lon = 'longitude',
4                          color = 'price_per_unit',
5                          width = 1000,
6                          height = 800,
7                          hover_data = ['price_per_unit'],
8                          zoom = 12,
9                          title = 'Estate Locations')
10
11 # stamen-terrain, open-street-map, carto-positron
12 fig.update_layout(mapbox_style = 'carto-positron')
13 fig.show()
```





In [6]:

```
1 # Create 3D scatter plot
2 fig = px.scatter_3d(re_data,
3                     x = 'latitude',
4                     y = 'longitude',
5                     z = 'price_per_unit',
6                     width = 800,
7                     height = 600,
8                     title = '3D view of Estate Locations')
9 fig.update_traces(marker = {'size': 4, 'line': {'width': 2, 'color': 'DarkSlateGrey'}},
10                  selector = {'mode': 'markers'})
11 fig.show()
```

In [7]:



```
1 # center = [24.98298, 121.54024]
2 # estate_locations = folium.Map(location = center, zoom_start = 13)
3 # for i, j in re_data.iterrows():
4 #     location = [j['latitude'], j['longitude']]
5 #     folium.CircleMarker(location, radius = 2).add_to(estate_locations)
6 # estate_locations
```



## Price distribution

```
In [8]: 1 display(re_data[['price_per_unit']].describe().transpose())
2 plt.figure(figsize = [10, 3])
3 plt.subplot(1, 2, 1)
4 hist_plot(re_data, 'price_per_unit', 'Distribution of Prices', 'Price per Unit', 'Frequency', 50, False)
5 plt.subplot(1, 2, 2)
6 box_plot(re_data, 'price_per_unit', None, 'Distribution of Prices', 'Price per Unit')
```

	count	mean	std	min	25%	50%	75%	max
price_per_unit	414.0	37.980193	13.606488	7.6	27.7	38.45	46.6	117.5



The summary statistics table, histogram and the boxplot, it is evident there are outlier values in the price and these needs to be sieved out.

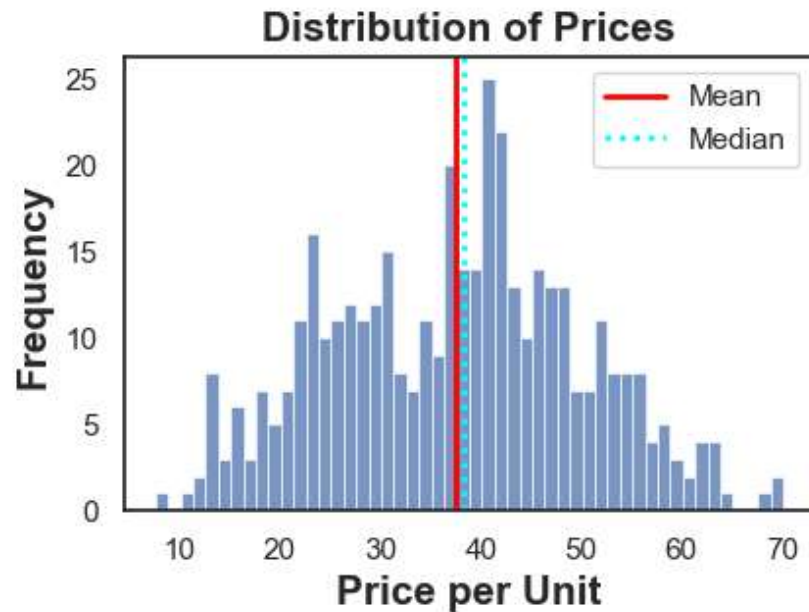
## Removing outlier values

The lower bound will be defined at 0%, and the upper bound at 99%.

In [9]:


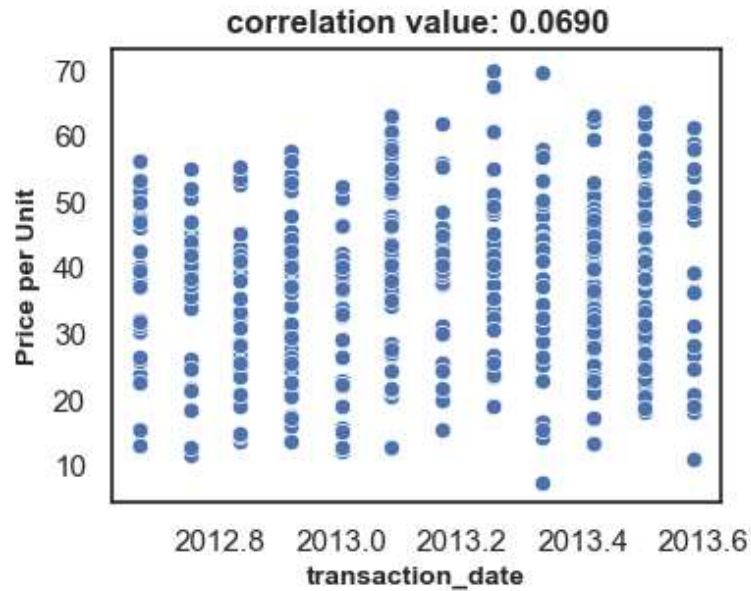
```
1 lower_bound = re_data['price_per_unit'].quantile(0)
2 upper_bound = re_data['price_per_unit'].quantile(0.99)
3 mark_bound_price = re_data['price_per_unit'].between(lower_bound, upper_bound)
4 re_data = re_data[mark_bound_price].reset_index().drop(columns = {'index'})
5
6 display(re_data[['price_per_unit']].describe().transpose())
7 plt.figure(figsize = [10, 3])
8 plt.subplot(1, 2, 1)
9 hist_plot(re_data, 'price_per_unit', 'Distribution of Prices', 'Price per Unit', 'Frequency', 50, False)
10 plt.subplot(1, 2, 2)
11 box_plot(re_data, 'price_per_unit', None, 'Distribution of Prices', 'Price per Unit')
```


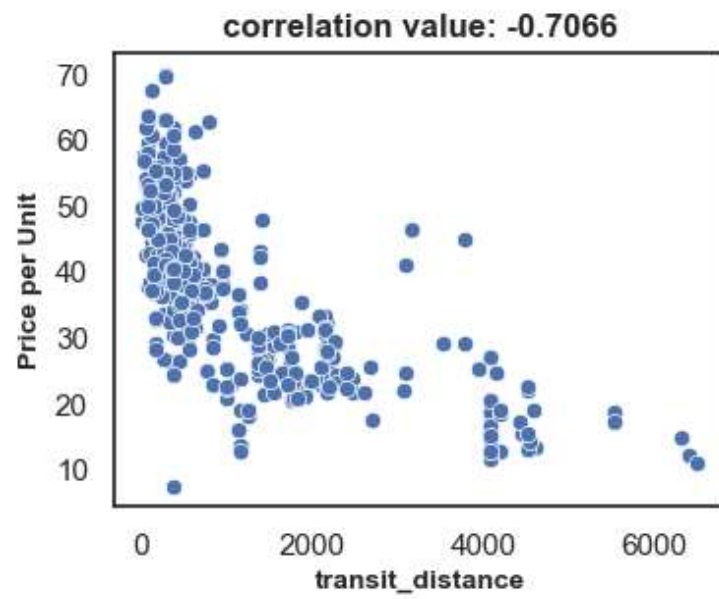
	count	mean	std	min	25%	50%	75%	max
price_per_unit	409.0	37.421516	12.565902	7.6	27.3	38.3	46.1	70.1


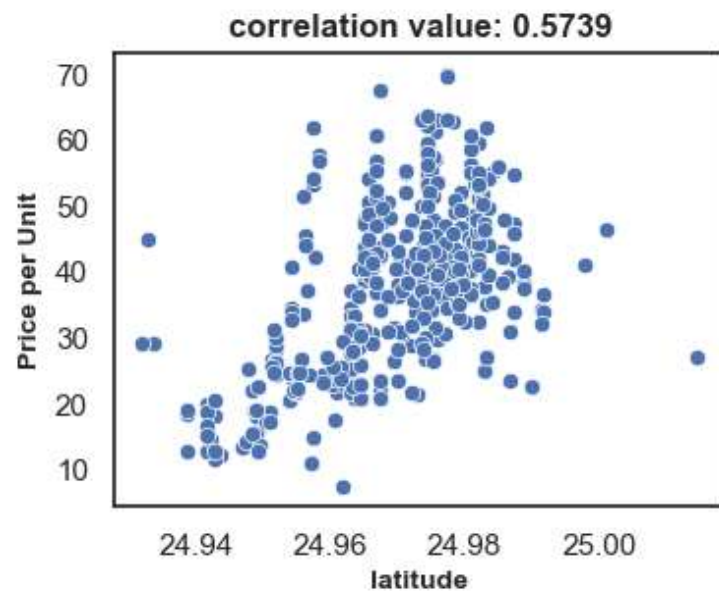


## Correlation between price and other features in the table

```
In [10]: 1 for col in re_data.columns[:-1]:
2         plt.figure(figsize = [4, 3])
3         correlation = re_data['price_per_unit'].corr(re_data[col])
4         scatter_plot(re_data, col, 'price_per_unit', col, 'Price per Unit', f'correlation value: {correlation:.4f}')
```








In [11]:

```
1 mask = np.triu(np.ones_like(re_data.corr()))
2 sb.heatmap(re_data.corr(), cmap = 'YlGnBu', annot = True, square = True, mask = mask, cbar = False);
```



The correlation values and the scatter plots above gives a clear evidence that house age , transit distance , local convenience stores , latitude and longitude all affect the price. Transaction date will be ignored since it has very negligible correlation with the price.

Checking the cardinality of local convenience stores:

```
In [12]: 1 print(f"Local convenience stores have {re_data['local_convenience_stores'].nunique()} unique values as listed below:")
2 display(re_data[['local_convenience_stores']].value_counts().sort_values())
```

Local convenience stores have 11 unique values as listed below:

local\_convenience\_stores

10	10
9	23
2	24
8	30
4	31
7	31
6	35
1	45
3	46
0	67
5	67

dtype: int64

## Separate labels, features, split and normalize data

```
In [13]: 1 # Separate features and Label
2 target_vector = re_data.columns[-1]
3 features = re_data.columns[1:-1]
4 X, y = re_data[features].values, re_data[target_vector].values
5
6 # split data: 70% for training and 30% for testing
7 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30, random_state = 0)
8 print(f'Training set: {X_train.shape[0]}\nTest set:{X_test.shape[0]}')
```

Training set: 286

Test set:123

# Build regression model

```
In [14]: 1 # Transform numeric columns
2 transformer = Pipeline(steps = [('scaler', StandardScaler())])
3 #numeric_features = ['house_age', 'transit_distance', 'Latitude', 'Longitude']
4 numeric_features = [0, 1, 2, 3, 4]
5 preprocess_data = ColumnTransformer(transformers = [('num', transformer, numeric_features)])
```

```
In [15]: 1 # Create a pipeline for Linear regression
2 model_reg = make_pipeline(transformer, preprocess_data, LinearRegression())
3 model_reg.fit(X_train, y_train)
```

```
Out[15]: Pipeline(steps=[('pipeline', Pipeline(steps=[('scaler', StandardScaler())])),
                        ('columntransformer',
                         ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('scaler',
                                                         StandardScaler())])),
                                                         [0, 1, 2, 3, 4])])),
                        ('linearregression', LinearRegression())])
```

```
In [16]: 1 # Pipeline form Ridge model
2 model_ridge = make_pipeline(transformer, preprocess_data, Ridge())
3 model_ridge.fit(X_train, y_train)
```

```
Out[16]: Pipeline(steps=[('pipeline', Pipeline(steps=[('scaler', StandardScaler())])),
                        ('columntransformer',
                         ColumnTransformer(transformers=[('num',
                                                         Pipeline(steps=[('scaler',
                                                         StandardScaler())])),
                                                         [0, 1, 2, 3, 4])])),
                        ('ridge', Ridge())])
```



```
In [17]: 1 # Pipeline for Random forest regressor
2 model_forest = make_pipeline(transformer, preprocess_data, RandomForestRegressor())
3 model_forest.fit(X_train, y_train)
```

```
Out[17]: Pipeline(steps=[('pipeline', Pipeline(steps=[('scaler', StandardScaler()))],
              ('columntransformer',
               ColumnTransformer(transformers=[('num',
                                               Pipeline(steps=[('scaler',
                                                                StandardScaler()))],
                                                                [0, 1, 2, 3, 4]))]),
              ('randomforestregressor', RandomForestRegressor())])
```

```
In [18]: 1 # Pipeline for GB Boost
2 model_gb_boost = make_pipeline(transformer, preprocess_data, GradientBoostingRegressor())
3 model_gb_boost.fit(X_train, y_train)
```

```
Out[18]: Pipeline(steps=[('pipeline', Pipeline(steps=[('scaler', StandardScaler()))],
              ('columntransformer',
               ColumnTransformer(transformers=[('num',
                                               Pipeline(steps=[('scaler',
                                                                StandardScaler()))],
                                                                [0, 1, 2, 3, 4]))]),
              ('gradientboostingregressor', GradientBoostingRegressor())])
```

# Evaluate Models

In [19]:

```
1  # Make predictions
2  reg_predictions = model_reg.predict(X_test)
3  ridge_predictions = model_ridge.predict(X_test)
4  forest_predictions = model_forest.predict(X_test)
5  gb_boost_predictions = model_gb_boost.predict(X_test)
6
7  # Get evaluation metrics
8  # Linear Regression
9  reg_mse = mean_squared_error(y_test, reg_predictions)
10 reg_rmse = np.sqrt(reg_mse)
11 reg_r2 = r2_score(y_test, reg_predictions)
12 print(f'Linear Regression results:\nMSE: {reg_mse}\nRMSE: {reg_rmse}\nR2_SCORE: {reg_r2}\n')
13
14 # Ridge
15 ridge_mse = mean_squared_error(y_test, ridge_predictions)
16 ridge_rmse = np.sqrt(ridge_mse)
17 ridge_r2 = r2_score(y_test, reg_predictions)
18 print(f'Ridge Regression results:\nMSE: {ridge_mse}\nRMSE: {ridge_rmse}\nR2_SCORE: {ridge_r2}\n')
19
20 # Random Forest
21 forest_mse = mean_squared_error(y_test, forest_predictions)
22 forest_rmse = np.sqrt(forest_mse)
23 forest_r2 = r2_score(y_test, forest_predictions)
24 print(f'Random Forest Regression results:\nMSE: {forest_mse}\nRMSE: {forest_rmse}\nR2_SCORE: {forest_r2}\n')
25
26 # GB Boost
27 gb_boost_mse = mean_squared_error(y_test, gb_boost_predictions)
28 gb_boost_rmse = np.sqrt(gb_boost_mse)
29 gb_boost_r2 = r2_score(y_test, gb_boost_predictions)
30 print(f'Random Forest Regression results:\nMSE: {gb_boost_mse}\nRMSE: {gb_boost_rmse}\nR2_SCORE: {gb_boost_r2}')
```

Linear Regression results:  
MSE: 43.01576372402771  
RMSE: 6.558640386850594  
R2\_SCORE: 0.6781479656322735

Ridge Regression results:  
MSE: 43.00464761992758  
RMSE: 6.5577928924240645  
R2\_SCORE: 0.6781479656322735

Random Forest Regression results:  
MSE: 26.975294438868193  
RMSE: 5.193774584911074  
R2\_SCORE: 0.7981657736331546

Random Forest Regression results:  
MSE: 26.323017673662722  
RMSE: 5.130596229841394  
R2\_SCORE: 0.8030462310672954

## Use trained models to predict price

```
In [20]: 1 data = {'house_age': [16.2, 13.6],  
2           'transit_distance': [289.3248, 4082.015],  
3           'local_convenience_stores': [5, 0],  
4           'latitude': [24.98203, 24.94155],  
5           'longitude': [121.54348, 121.50381]}  
6 predict_data = pd.DataFrame.from_dict(data)  
7 predict_data
```

```
Out[20]:
```

	house_age	transit_distance	local_convenience_stores	latitude	longitude
0	16.2	289.3248	5	24.98203	121.54348
1	13.6	4082.0150	0	24.94155	121.50381

In [21]:

```
1 # Linear Regression
2 linear_price = model_reg.predict(predict_data)
3 print(f'Linear Regression predicted prices:')
4 for predicted_price in linear_price:
5     print(predicted_price.round(2))
6
7 # Ridge
8 ridge_price = model_ridge.predict(predict_data)
9 print(f'\nRidge predicted prices:')
10 for predicted_price in ridge_price:
11     print(predicted_price.round(2))
12
13 # Random Forest
14 forest_price = model_forest.predict(predict_data)
15 print(f'\nRandom Forest Regression predicted prices:')
16 for predicted_price in forest_price:
17     print(predicted_price.round(2))
18
19 # Gradient Boosting
20 gb_price = model_gb_boost.predict(predict_data)
21 print(f'\nGradient Boosting Regressor predicted prices:')
22 for predicted_price in gb_price:
23     print(predicted_price.round(2))
```

Linear Regression predicted prices:

45.59

15.26

Ridge predicted prices:

45.58

15.3

Random Forest Regression predicted prices:

49.27

15.82

Gradient Boosting Regressor predicted prices:

49.33

16.73

## Save models

```
In [22]: 1 reg = 'linear_regression_model.pkl'
          2 ridge = 'ridge_regression_model.pkl'
          3 forest = 'random_forest_regression_model.pkl'
          4 gb = 'gradient_boosting_regressor_model.pkl'
          5
          6 joblib.dump(model_reg, reg)
          7 joblib.dump(model_ridge, ridge)
          8 joblib.dump(model_forest, forest)
          9 joblib.dump(model_gb_boost, gb)
```

```
Out[22]: ['gradient_boosting_regressor_model.pkl']
```